

## Ranking

The `rai ranking` command is now active. The `ranking` command calculates the fastest sum of both Layer Runtimes for all valid runs for each student and then ranks each student from fastest to slowest. *Rankings* are updated immediately with each valid run. A valid run requires no profiling and a batch size of 10,000.

Below is an example of *ranking* results. Notice that your netid and “-->” indicate your rank and time.

Ranking:

YOU	RANK	FASTEST (MS)
	0	1.040769s
	1	1.051001s
	2	1.06379s
	3	1.167731s
	4	1.188924s
aschuh3 -->	5	1.204951s

Note: We began collecting *ranking* statistics on 4/15/2021 at 11:38 AM. Prior runs will not be included in *ranking*.

## Exclusive Queue

We currently have 4 servers, each with 2 GPUs, which each run up to 6 simultaneous jobs. This allows up to a total of 24 simultaneous jobs which is perfect for handling a large number of jobs efficiently. However, profiling results are impacted significantly when other programs compete for resources.

To address this issue, the exclusive queue is now active and should be used for profiling only. Three of the servers will handle the default `rai_amd64_ece408` queue, allowing 18 simultaneous runs. The fourth server will now handle the `rai_amd64_ece408_exclusive` queue which will allow 2 simultaneous runs (1 on each GPU). The `rai` api allows you to specify the queue with `-q <<queue name>>`. Note, the default queue is `rai_amd64_ece408` so the `-q` option is only required when using the `rai_amd64_ece408_exclusive` queue.

Example:

```
rai -p <<address of your project code>> -q rai_amd64_ece408_exclusive
```

There will be times when submitted jobs may not run immediately because of a large queue. This is particularly true when using the `rai_amd64_ece408_exclusive` queue. Below are a few tips when working with a large queue backlog.

- You can check the current queue size by adding “queued” to the end of your `rai` command. The `rai` command must still include `-p <<address of your project code>>` and will default to returning the `rai_amd64_ece408` queue size. Include `-q rai_amd64_ece408_exclusive` to see the current `rai_amd64_ece408_exclusive` queue size.
- After your job is submitted (waiting for server), you can kill the `rai` command and even turn off your computer. A server will handle your queued job in its time and post the results to the database. Later you can run the `rai history` command. If the job has run it will be included in your history request and the results can be downloaded through the associated URL.

Note: As needed we may dynamically move other machines to the `rai_amd64_ece408_exclusive` queue.